

An introduction to RL and deep RL (2)

BABAK BADNAVA

MULTI-AGENT SYSTEMS LAB, IUST

This slide is mainly taken from silver's slides



مرکز تحقیقات هوش پارت

Deep neural networks for supervised tasks

- Data: sample, label
- Using deep neural networks:
- h_i could be any kind of function
- We have:
 - A sample (x)
 - Its label (y^*)
 - Network's prediction (y)

$$x \xrightarrow{w_1} h_1 \xrightarrow{w_2} \dots \xrightarrow{w_n} h_n \xrightarrow{w_{n+1}} y$$

Deep neural networks for supervised tasks (2)

- Define a loss function: $x \xrightarrow{w_1} h_1 \xrightarrow{w_2} \dots \xrightarrow{w_n} h_n \xrightarrow{w_{n+1}} y \longrightarrow l(y)$
- Loss function could be:
 - Mean-squared error: $l(y) = \|y^* - y\|^2$
 - Log likelihood: $l(y) = \log P[y^*|x]$
 -
- Minimize loss
- Using different optimization algorithms

Overview of approximation methods

- Value estimation
 - Use a neural network to approximate value function
 - Define a policy function on top of it
- Policy gradient
 - Use a neural network to approximate policy function directly

Let's apply it to SARSA

- Action value function:
 - $Q_{\pi}(s, a) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s, A_t = a]$
- Represent it by Q-network with weights w
 - $Q_{\pi}(s, a) \approx Q(s, a, w)$
- Define loss function:

$$\mathcal{L}(w) = \mathbb{E} \left[\left(\underbrace{r + \gamma Q(s', a', w)}_{\text{target}} - Q(s, a, w) \right)^2 \right]$$

How about Q-learning?

- Loss function:

$$\mathcal{L}(w) = \mathbb{E} \left[\left(\underbrace{r + \gamma \max_{a'} Q(s', a', w)}_{\text{target}} - Q(s, a, w) \right)^2 \right]$$

- Optimize this using SGD, using $\frac{\partial \mathcal{L}(w)}{\partial w}$

$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a', w) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$

Summary

- Use this network to estimate value function instead of a Q-table



- Nothing else have been changed
- Does it work well?

Compare

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Algorithm from Sutton 2018

Compare (2)

- Initialize weights
- Loop for each episode:
 - Initialize s
 - Choose a from s using policy derived from $Q(s, a, w)$
 - Loop for each step of episode:
 - Take action a , observe r, s'
 - Choose a' from s' using policy derived from $Q(s', a', w)$
 - Compute gradient of loss and update the network
 - $s \leftarrow s'; a \leftarrow a'$
 - Until s is terminal

Does it work well?

- Oscillates or diverges
- Why?
 - Data is sequential
 - Successive samples are correlated
 - Policy changes rapidly with slight change to Q-values
 - Distribution of data will change
 - Scale of rewards and Q-values is unknown

What to do?

- Use experience replay
 - Break correlation between data
- Freeze target Q-network
 - Avoid oscillations
- Clip rewards or normalize network adaptively

Stable Deep RL: experience replay

- To remove correlations, build data-set from agent's own experience
- Take action a_t according to epsilon-greedy policy
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory
- Sample random mini-batch of transitions (s, a, r, s') from replay memory
- Optimise MSE

$$\mathcal{L}(w) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[\left(r + \gamma \max_{a'} Q(s', a', w) - Q(s, a, w) \right)^2 \right]$$


Stable Deep RL: fixed target Q-network

- Use target Q-network with fixed parameter
- Choose action based on the target Q-network
- Compute Q-learning target w.r.t old, fixed parameter
- Optimize MSE
- Periodically update fixed parameters $w^- \leftarrow w$

$$\mathcal{L}(w) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[\left(r + \gamma \max_{a'} Q(s', a', w^-) - Q(s, a, w) \right)^2 \right]$$

DQN architecture in Atari

- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s is stack of raw pixels from last 4 frames
- Output is $Q(s, a)$ for 18 joystick/button positions
- Reward is change in score for that step



Policy gradient

- Represent policy by deep network $a = \pi(s, u)$ with weights u
- Define objective function as total discounted reward

$$J(u) = \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots]$$

- Optimise objective end-to-end by SGD
 - Adjust policy parameters u to achieve more reward

$$\begin{aligned} \frac{\partial J(u)}{\partial u} &= \mathbb{E}_s \left[\frac{\partial Q^\pi(s, a)}{\partial u} \right] \\ &= \mathbb{E}_s \left[\frac{\partial Q^\pi(s, a)}{\partial a} \frac{\partial \pi(s, u)}{\partial u} \right] \end{aligned}$$

Actor-Critic method

- Actor is a policy $\pi(s, u)$ $s \xrightarrow{u_1} \dots \xrightarrow{u_n} a$

- Critic is value function $Q(s, a, w)$ with parameter w

$$s, a \xrightarrow{w_1} \dots \xrightarrow{w_n} Q$$

- Critics provide loss function for actor

$$s \xrightarrow{u_1} \dots \xrightarrow{u_n} a \xrightarrow{w_1} \dots \xrightarrow{w_n} Q$$

- Gradient back propagates from critic into actor

$$\frac{\partial a}{\partial u} \longleftarrow \dots \longleftarrow \frac{\partial Q}{\partial a} \longleftarrow \dots \longleftarrow$$

Actor-Critic: Learning rules

- Critic estimates value of current policy by Q-learning

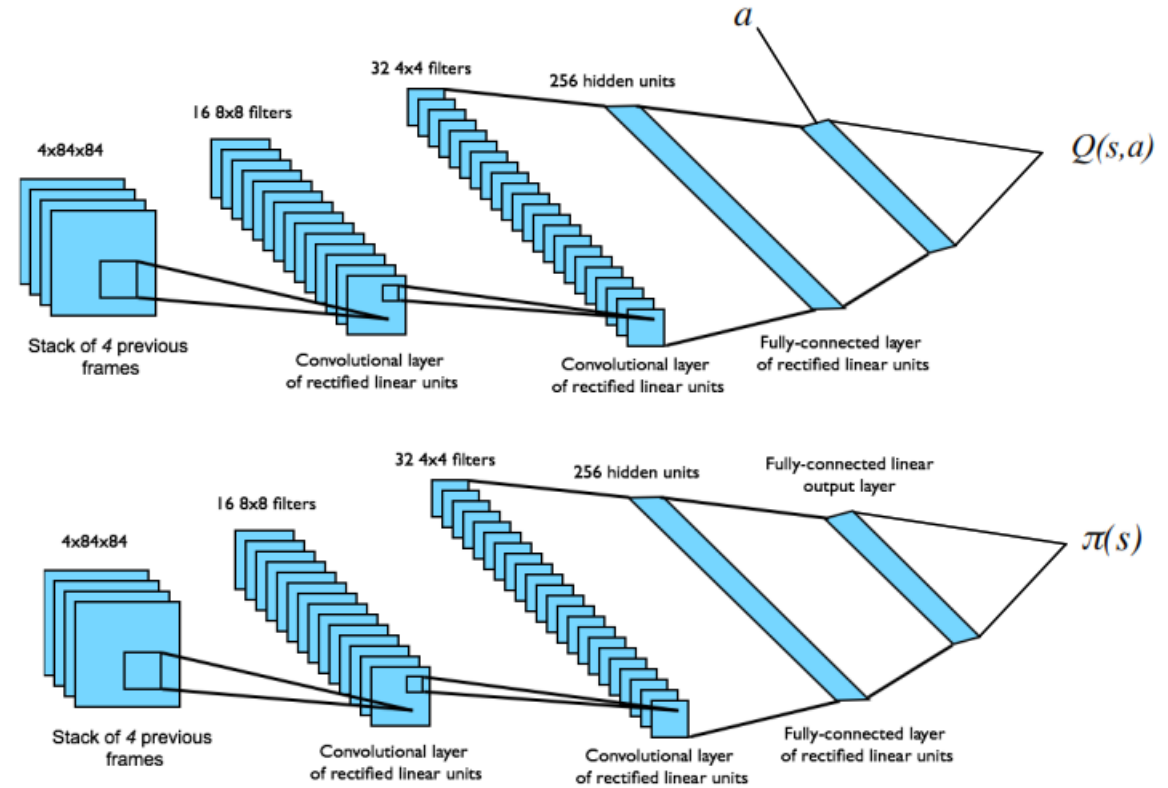
$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E} \left[\left(r + \gamma Q(s', \pi(s'), w) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$

- Actor updates policy in direction that improves Q

$$\frac{\partial J(u)}{\partial u} = \mathbb{E}_s \left[\frac{\partial Q(s, a, w)}{\partial a} \frac{\partial \pi(s, u)}{\partial u} \right]$$

Architecture and tips


- Use experience reply
- Freeze target network



How to improve?

- Run multiple simulation simultaneously instead of reply memory
 - Try this one as a practice
 - And think about its advantageous
- Use Transfer Learning
- ...

A simple python code



References

- David Silver's slide
- An introduction to Reinforcement Learning by Sutton 2nd edition
- Berkeley's deep RL boot camp materials available at [here](#)

Questions?