

# An introduction to RL and deep RL

BABAK BADNAVA

MULTI-AGENT SYSTEMS LAB, IUST



# Different type of machine learning

- Supervised
  - Labeled data
- Unsupervised
  - Unlabeled data
- Semi-supervised
  - A small amount of labelled data with a large amount of unlabelled data
- Reinforcement learning
  - Data have no label, but there is a feedback.

# Reinforcement learning

- Agent-oriented learning
- Learning by interacting with an environment to achieve a goal
- Learning by trial and error, with only delayed evaluative feedback (reward)
- The kind of machine learning most like natural learning
- Learning that can tell for itself when it is right or wrong

## Reinforcement learning (2)

- An agent observes the environment
- Makes a decision
- And gets a feedback


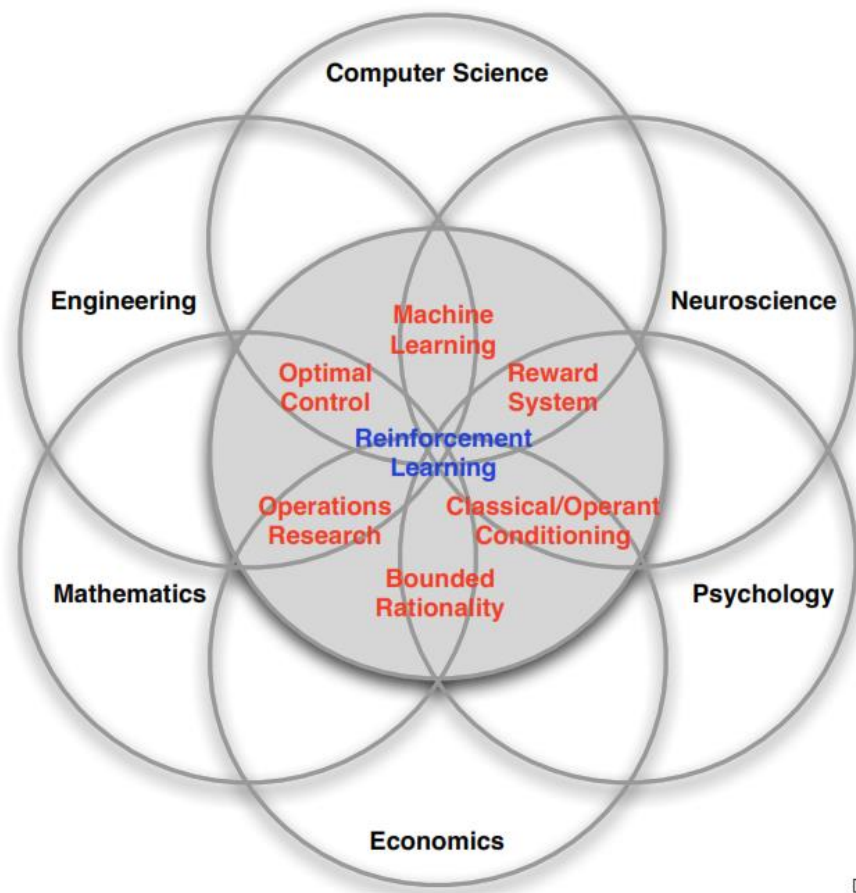


Image taken from Wikipedia.org

# Where it came from?



David Silver 2015




# Applications

- Robotics
- Personalized web services
- Neuroscience
- Autonomous cars
- ...

# How to formulate?

- Use Markov decision process (MDP) to formulate.
- MDP is a tuple:  $(S, A, P, R, \gamma)$
- $S$  is a set of all possible states
- $A$  is a set of all actions
- $P$  is a transition function
- $R$  is a reward function



# Example: autonomous vehicle

- States: a set of all possible value for all sensors, speed, amount of available gas, ...
- Actions: Pedals, Steering Wheels, ...
- Reward: amount of money for an autonomous cab, distance to the destination, .....
- Transition function: natural rules + human-made rules
  - Gravity
  - Driving rules
  - ...



# Practice

- Formulate an MDP for an online shopping website who wants to increase its profit by giving its customers a discount.
- Formulate an MDP for a classification task like MNIST
- Formulate an MDP for a manufacturing robot

# Policy function

- Map a state  $s$  to action  $S \times A \rightarrow [0,1]$
- Can be deterministic or stochastic:
  - Deterministic:  $\pi(s) = a$
  - Stochastic:  $\pi(a|s) = P_{\pi} [A = a|S = s]$
- Example:
  - Car changing gear.

# Value function

- How good is being in a state?
  - $V_{\pi}(s) = \mathbb{E}[G_t | S_t = s]$
  - $G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$
- Action-value function:
  - $Q_{\pi}(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$
  - $V_{\pi}(s) = \sum_{a \in A} Q_{\pi}(s, a) \pi(a|s)$

# Optimal value and policy

*Optimal value functions:*

- $V_*(s) = \max_{\pi} V_{\pi}(s)$
- $Q_*(s, a) = \max_{\pi} Q_{\pi}(s, a)$
- Optimal policy:
  - $\pi_* = \arg \max_{\pi} V_{\pi}(s)$

# Different type of MDP and how to solve them?

- Model-based:
  - We know both the transition function and reward function.
  - Dynamic programming.
- Model-free:
  - No transition function.
  - No reward function.

# Model-free

- Value estimation
  - Estimate the value function
  - Use a predefined soft policy function
    - Epsilon greedy
    - Boltzmann policy
- Policy approximation
  - Directly approximate the optimal policy function

# Value estimation

- Update value function till convergence
  - $V(s_t) \leftarrow (1 - \alpha)V(s_t) + \alpha G_t$
  - $V(s_t) \leftarrow V(s_t) + \alpha(R_{t+1} + \gamma V(s_{t+1}) - V(s_t))$
- We can do this for action-value function too:
  - $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$
  - We call this SARSA
  - This is an on-policy method
- Q-learning: an off-policy method:
  - $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_{a \in A} Q(s_{t+1}, a) - Q(s_t, a_t))$

## Soft policies

- Epsilon greedy:

- $\pi(a|s) = \begin{cases} \varepsilon & \text{if } a = \arg \max_a Q(s, a) \\ \frac{1-\varepsilon}{|A|-1} & \text{otherwise} \end{cases}$

- Boltzmann (softmax):

- $\pi(a|s) = \frac{e^{-Q(s,a)/\tau}}{\sum_{a \in A} e^{-Q(s,a)/\tau}}$



# Let's put it all together: SARSA

Sarsa (on-policy TD control) for estimating  $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

    until  $S$  is terminal

Algorithm from Sutton 2018

# Let's put it all together: Q-learning

Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$


$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

    until  $S$  is terminal

Algorithm from Sutton 2018

# A sample python code



Sutton 2018

# Practice

- Try to implement SARSA and Q-learning for a continuous environment like Mountain Car.

# References

- An introduction to Reinforcement Learning by Sutton 2<sup>nd</sup> edition